

Summary

Machine Learning Operations Engineer with 5 years' experience, combining ML and software engineering. Led cross-functional teams and redesigned core platform architecture, cutting costs by 40% and managing Kubernetes infrastructure.

Skills

- Programming Languages: Python
- Big data tools: Elasticsearch, Postgres, BigQuery, S3
- Machine learning: NLP, PyTorch, Scikit-learn, LLMs, Generative AI, MCP
- Frameworks : Fast API, Pandas
- Devops : Kubeflow, Flyte, Docker, ArgoCD, SQS, GCP, AWS, Lambda, Kubernetes, Grafana, Prometheus, Keda, Karpenter

Experience

Pixis

Machine Learning Operations Engineer

Bengaluru, India.

December 2024 – Present

- Managed AWS Kubernetes infrastructure using Terraform; streamlined deployments with GitHub Actions and Argo CD.
- Deployed scalable LLM & GPU models for ad platforms via Flyte with sub-120ms; migrated pipelines from Kubeflow.
- Optimized Docker image caching, accelerating feature calculation by 2× and reducing cold start times by 60%.
- Configured AWS Lambdas to process 2M+ daily payloads, using Poetry package management and AWS CodeArtifact.
- Architected and deployed MCP server to enable experimentation with AI agents for the data science team.
- Monitored ML workflows with Grafana and Prometheus, achieving 99.99% uptime; researched Spark Operator for feature calculation.
- Review code changes and model building from data scientists and give feedback.

Stylumia (founded by ex-COO of Myntra, acquired by Walmart)

Machine Learning Engineer

Bengaluru, India.

August 2020 – December 2024

- Developed a serverless model serving tool, utilizing Knative on GCP Kubernetes to host ML model API servers with automatic scaling—seamlessly scaling down to zero during inactivity and rapidly scaling up in response to new requests.
- Enhanced the ML platform's capacity to handle inference workloads for 1 million products in a single day on GKE.
- Achieved a 40% reduction in infrastructure costs by optimizing resource utilization.
- Led a cross-functional team to build and deploy an automated NLP attribute extraction platform using GCP Vertex AI, MLflow and consensus of LLMs (OpenAI, Anthropic, Vertex AI) achieving over 90% accuracy and 95% coverage.
- Built robust ETL pipelines using Elasticsearch, BigQuery, & SQS, reducing data latency for ML model inference by 60%.
- Enhanced retail **product matching** by fine-tuning open-source Flag Embedding with triplet loss; achieved 99% precision for fashion and 93% for home & hardware.

Crisil Global Research and Analytics (S&P)

Data Science Intern

Mumbai, India.

June 2019 – July 2019

- Developed news ml pipeline with sentiment and trend analysis, improving financial equity research report generation speed by 60% in the product, SMART.

Crowd Analytix Pvt. Ltd.

Data Science Intern

Bengaluru, India.

June 2018 – July 2018

- Built classification models for e-commerce clients(Walmart, The Home Depot) for improving customer experience on websites by recommending relevant products according to product attributes with accuracy of 90%.

Honors & Awards

- **Superstar Award (Twice) and Value Champion Award** - Recognized at Stylumia for exceptional performance, innovation, and exemplifying company values.
- **ZS Machine Learning Challenge 2019** - Achieved 289 Rank out of 6000 participants in the annual Hackathon hosted by ZS on interviewBit platform.

Education

Maharaja Agrasen Institute of Technology, GGSIPU

B.Tech, Electronics and Communication, 8.1 C.G.P.A

Delhi, India

August 2016 - August 2020